

Guidelines



Judging Criteria

Dimension 1: Impact Potential & Innovation

How much would this matter for AI safety if it worked? How innovative is it?
For scores of 4-5: is this actually new to the field, or replicating recent work?

Score	Description
1	Negligible. No clear problem addressed, or no meaningful novelty.
2	Limited. Addresses a real problem but with a generic or well-trodden approach. Incremental at best.
3	Moderate. Clear problem with a reasonable approach; some novelty in framing or method beyond routine application of existing tools.
4	Significant. Important problem with an original approach, or identifies a neglected problem area. A valuable contribution others could build on.
5	Exceptional. Tackles a critical AI safety problem with a genuinely novel approach, or opens a new research direction. Clear theory of change. You'd be excited to share this with researchers in the area.

Dimension 2: Execution Quality

How sound are methodology, implementation, and findings?

Score	Description
1	Seriously flawed. Methodology broken, results uninterpretable, or implementation doesn't work.
2	Weak. Approach has significant gaps: missing validation, flawed experimental design, or incomplete implementation.
3	Competent. Technically solid given the short duration. Methodology makes sense, results are interpretable, limitations acknowledged, work builds toward clear conclusions.
4	Strong. Thorough methodology with convincing validation. Results clearly support conclusions. Immediately useful for future work.
5	Exceptional. Ambitious scope executed rigorously. Surprising findings, novel methods, or unusually robust validation.

Dimension 3: Presentation & Clarity

How clearly are work, findings, and impact potential communicated?

Score	Description
1	Incomprehensible. Cannot determine what the project is actually claiming or doing.
2	Hard to follow. Key information buried, missing, or diluted by excessive length. Significant effort to extract main points.
3	Clear enough. Can understand the problem, approach, and results without undue effort. Core content clearly present: problem, method,

	findings, limitations.
4	Well presented. Easy to follow, well-structured, appropriate level of detail. Target audience would get it quickly.
5	Exceptionally clear. A pleasure to read. Complex ideas made accessible. Could serve as a model for how to present this type of work.

Submission Requirements

A complete submission includes:

1. A research report in PDF format using the official [submission template](#)
2. A project title and brief abstract (150 words max)
3. Author names and affiliations for all team members
4. An indication of which challenge track your project addresses

Report structure (recommended):

There is no hard page limit. Most winning projects are 4–8 pages. The report should include:

- Introduction: What problem did you address? Why does it matter for AI control?
- Related Work: What existing work does your project build on?
- Methodology: What did you build or test? Describe your approach in enough detail for someone to replicate it
- Results: What did you find? Include quantitative results where possible
- Discussion: What are the implications? What are the limitations? What would you do with more time?
- References: Cite relevant prior work

Important notes:

- You can submit as an individual or as a team
- You can build on existing work, but you **MUST** clearly identify what is **NEW** work done during the hackathon

- hackathon
- Your PDF is not editable or replaceable once submitted
- Submit through the official submission form on the hackathon page
- If you run into submission issues, DM Kamil on Discord or email sprints@apartresearch.com

Submission Template ⇒ [Link](#)

Frequently Asked Questions

Getting Started

Q: How does the hackathon work?

A: Sign up, join the Discord server, form or join a team (or work solo), pick a direction, build your project over three days, and submit a research report (PDF) by the deadline. Talks and Q&A sessions run throughout the event.

Q: How long is the hackathon?

A: Three days, March 20-22, 2026. Submissions are due end of day Sunday, the last day of the event. The full schedule with speaker times is on the Schedule tab.

Q: Can I participate remotely?

A: Yes. The hackathon is fully virtual. All talks, collaboration, and submissions happen online through Discord and Zoom.

Q: How do teams work?

A: Teams form before or during the hackathon. Check the team-forming channels on Discord to find collaborators. Solo participation is fine. We recommend teams of up to 5, but larger groups are allowed.

Q: Do tracks affect scoring?

A: No. Tracks help you find a focus area and relevant resources, but they do not affect how your project is scored. You can work on something that spans multiple tracks.

Q: How do I get Lambda cloud credits?

A: After signing up you will receive an email with the Lambda credits link. Credits are per team, not per person. You need to create a Lambda account and verify with a credit card (\$10 temporary hold, refunded). Apply before Friday if the hackathon runs over the weekend, since Lambda processes applications during business hours only.

applications during business hours only.

Q: Are the talks recorded?

A: Yes. Recordings are shared on Discord after the event.

Submissions

Q: What do I submit?

A: A research report in PDF format using the [submission template](#). Think of it as a mini research paper documenting your problem, approach, results, and implications. Not a product demo.

Q: Which submission template should I use?

A: Always use the one linked on the Guidelines tab of the hackathon website. The template in the acceptance email may be an older or minimal version.

Q: Will I get a confirmation after submitting?

A: Yes. You will receive a confirmation email with your project title shortly after submitting. If you do not receive one, contact sprints@apartresearch.com.

Q: My project doesn't show up on the website after submitting.

A: Submissions are reviewed and published manually. It can take up to 12 hours for your project to appear on the website. If it is still missing after that, email sprints@apartresearch.com.

Q: I submitted a duplicate by accident.

A: DM Kamil on Discord with your project title. He will keep the correct submission and delete any extras.

Q: Can I fix my PDF after submitting?

A: Your PDF is not editable or replaceable once submitted. Double-check your report before submitting. If there is a critical error, email sprints@apartresearch.com, but corrections are not guaranteed.

Q: Can I add team members after submitting?

A: Yes. You can update team members through the submission form on the website. If you need help, DM Kamil on Discord.

Q: I forgot to indicate my track in the submission.

A: No action needed. Tracks do not affect scoring.

Q: Can I submit an unfinished project?

A: Yes. Submitting something unfinished is always better than not submitting. Judges evaluate what you accomplished during the hackathon timeframe.

Q: Can I build on existing research?

A: Yes, but you must clearly identify what is new work

A: Yes, but you must clearly identify what is new work done during the hackathon. Projects with significant undisclosed prior work will be disqualified, even from top prize positions. If a judge or organizer contacts you about prior work concerns, respond promptly. Failure to respond may result in disqualification.

Judging & Results

Q: How does judging work?

A: Your project is assigned to 2-3 expert judges who review your PDF and score it on the three dimensions above. Judges typically have about one week after the event to complete reviews.

Q: Where is the scoring rubric?

A: Above on this page.

Q: Are individual judge scores shared?

A: No. The rubric is public, but individual scores are internal. Constructive feedback from judges is shared with participants without reviewer names.

Q: When will results be announced?

A: Typically 1-2 weeks after the judging deadline. Winners are contacted directly. All participants receive reviewer feedback by email.

Support

Q: How do I get help during the hackathon?

A: Three ways:

- Discord help-desk channel (tag @Kamil Alaa)
- DM Kamil on Discord
- Email: sprints@apartresearch.com

Speakers & Collaborators





Aryan Bhatt

Keynote Speaker

Senior Member of Technical Staff at Redwood Research, focused on AI control. Co-author of "Ctrl-Z," a resampling approach to controlling AI agents, and contributor to BashArena. Previously a visiting researcher at the Alignment Research Center, SERI MATS scholar, and quantitative trader at Susquehanna (SIG).



Buck Shlegeris

Judge and Co-organizer

CEO of Redwood Research, a nonprofit focused on mitigating risks from advanced AI. A leading voice on AI control: safely deploying AI systems that might be misaligned through monitoring, oversight, and containment.



Tyler Tracy

Speaker, Judge and Co-organizer

Experienced software engineer turned AI control researcher at Redwood Research, focused on developing technical safeguards for advanced AI systems with expertise in monitoring frameworks. Mentors at CBAI and Pivotal Research on control evaluations and protocol design.



Jason Hoelscher-Obermaier

Judge and Organizer

Director of Research at Apart Research. Ph.D. in physics from Oxford. Focuses on AI safety evaluations, interpretability, and alignment.



Kamil Alaa

Organizer

Operations at Apart Research, managing research sprints and hackathons.



Akshay Iyer

Judge and Co-organizer

CS and Entrepreneurship at Columbia University, IIT Bombay alum. Research experience in neuromorphic engineering and federated learning. Apart Research judge and internal organizer.



Rogan Inglis

Speaker and Judge



Jai Dhyani

Speaker and Judge

Senior Research Engineer on the Control team at the UK AI Security Institute (AISI). 7+ years of experience building and evaluating AI safety systems.

Builder of Luthien Proxy at Luthien Research, bringing Redwood-style AI control to real deployments. Co-author of RE-Bench (ICML 2025) with Elizabeth Barnes at METR.



Sebastian Farquhar
Speaker


Staff Research Scientist at Google DeepMind on the AGI Safety and Alignment Team, leading the Frontier Safety Loss of Control team. Senior Research Fellow at OATML Oxford.



Fabien Roger
Speaker

AI safety researcher and co-author of the foundational AI control paper (ICLR 2025). Former Member of Technical Staff at Redwood Research. His work spans alignment faking, chain-of-thought monitoring, and capability elicitation.





James Lucassen

Speaker and Judge

Researcher at Redwood Research working on AI control. His work spans AI benchmarking, LLM behavior analysis, and safety evaluations. Previously published on false belief detection in language models and AI performance prediction.



Ram Potham

Speaker and Judge


Astra Fellow at Redwood Research, focused on mitigating loss-of-control risk. His agent safety research was accepted for an oral presentation at the ICML 2025 Technical AI Governance workshop.



Nick Kuhn

Speaker and Judge

Postdoctoral researcher at the University of Oxford, collaborating with Dominic Joyce. PhD from Stanford University. Research focuses on moduli spaces in algebraic geometry. Previously at the University of Oslo and Max Planck Institute for Mathematics in Bonn.



Myles Heller

Speaker and Judge

Freelance programmer contributing to AI model safety evaluations at EquiStamp, designing benchmarks and evaluation tasks for AI agents. Volunteer developer at AISafety.info, building a RAG chatbot, and contributor to AI-Plans.com.



Justin Shenk

Speaker and Judge

Independent AI researcher and SPAR mentor based in Berlin, contracting at Redwood Research on model evals. Background in machine learning and software engineering.



Francis Rhys Ward

Judge

Researcher at Safe & Trusted AI and Future of Life Institute. Co-author of "CTRL-ALT-DECEIT" on sabotage evaluations (NeurIPS 2025 Spotlight), finding that sandbagging is much harder to detect than code sabotage.



Mia Hopman

Judge

Member of Technical Staff at Apollo



Anshul Khandelwal

Judge

Researcher at Redwood Research. His

Member of Technical Staff at Apollo Research, working on alignment, control, and security of large language models. Published on evaluating scheming propensity in LLM agents.

Researcher at Redwood Research. His work spans alignment drift, power-seeking evaluations, and white-box control methods.



Alex van Grootel

Judge

Fellow at the Future of Life Foundation with 10 years of experience applying AI to materials science and decision-making.

Previously a Product Manager at Microsoft (Fabric/AI copilots) and Data Scientist Team Lead at Citrine Informatics. MS from MIT.



Cameron Tice

Judge

Co-Director and Co-Founder of Geodesic Research, a technical AI safety organization at the University of Cambridge. Marshall Scholar and former Apart Research Fellow.





Theo Ryzhenkov

Judge

AI Safety Research Engineer and Founding Engineer at Palisade Research. NeurIPS 2025 author on detecting sandbagging in language models. SPAR, ARENA, and AISF alumni.



Pablo Bernabeu Perez

Judge

AI Safety Researcher and Research Engineer at the Barcelona Supercomputing Center. Co-author of "CoT Red-Handed" on stress-testing chain-of-thought monitoring at LASR Labs. Research Mentor at SPAR and Algoverse, with work on debate protocols for AI control accepted at NeurIPS 2025.



Akash Kundu

Judge

Research Collaborator at FAR.AI working on AI safety, LLM evaluations, and multilingual model safety. ICLR 2025 Oral presentation recipient and upcoming Cooperative AI Research Fellow.



Ramanpreet Singh Khinda

Judge

Staff Software Engineer at LinkedIn. IEEE Senior Member, Google I/O Codelab author, and experienced hackathon judge with expertise in mobile AI systems.



Hugo Delahousse

Judge

Software Engineer at Charcoal, an AI retrieval startup backed by investors from Cursor, Notion, and Front. Previously at Front for 5+ years building developer tools. 10 years of software engineering experience with a focus on AI automation and workflow security.



Victor Gillioz

Judge


Victor Gillioz is an AI safety researcher and MATS Fellow. His work spans training-time alignment of LLMs, including recontextualization for reward hacking mitigation and inoculation prompting, a technique for improving model alignment at inference time, under Alex Turner. His current work focuses on training scheming monitors for LLM oversight under Marius Hobbhahn, aiming to detect and contain misaligned model behavior.



Jord Nguyen

Judge

Research Manager at antoan.ai (AI Safety Vietnam) and Founder of the Hanoi AI



Kaushik Prabhakar

Judge

Research Fellow at Apart Research. His research spans AI safety and alignment

Vietnam), and Founder of the Hanoi AI Safety Network. Former Apart Research Fellow and Non-trivial facilitator.

research spans AI safety and alignment, with work on LLM evaluations. ARENA Cohort 4 alumni.



Helios Lyons

Judge

Embedded Software Engineer at Disguise and former Research Fellow at Apart Research. His AI safety work focuses on detecting sycophancy and dark patterns in language model



Nguyen Nhat Minh

Judge

AI Engineer at Mobifone IT Center and researcher at BKAI Lab, Hanoi University of Science and Technology. Published at FSE and TechDebt on vulnerability detection using graph neural networks

