
PROJECT TITLE¹

Author name 1
Affiliation

Author name 2
Affiliation

Author name 3
Affiliation

Author name 4
Affiliation

Author name 5
Affiliation

Author name 6
Affiliation

With
Apart Research

Abstract

Summarize your project in 150–250 words. A good abstract covers: the problem you addressed, why it matters, your approach, key results, and the main takeaway or implication.

How to use this template: Replace the *italicized guidance text* under each section with your content. The section structure is strong guidance but not rigid. If your project requires a different organization, feel free to adapt. Delete all guidance text including this info box before submitting. Make sure the project title and author information above are up to date.

How your submission will be evaluated: Your project will be judged on the quality of this written report. To understand what we look for, see our [Evaluation Rubric](#).

Recommended length: 4 pages excluding references and appendix.

¹ Research conducted at the [AI Control Hackathon](#), March 2026

1. Introduction

What problem are you addressing and why does it matter? Provide enough background for readers to understand your work. If relevant, briefly describe the threat model or failure mode you're targeting. Aspire to clearly list your most important contributions.

Our main contributions are:

1. *[First contribution — what new thing did you create, discover, or demonstrate?]*
2. *[Second contribution]*
3. *[Third contribution, if applicable]*

2. Related Work

What prior work is most relevant, and how does your work differ? Cite the most relevant papers, tools, or projects. Explain what gap your work addresses.

3. Methods

Describe your approach clearly enough that someone could replicate it. Include key design choices and justify them where relevant.

4. Results

Present your main findings with appropriate evidence. Use figures and tables where appropriate (see tips below). Distinguish between observations and interpretations.

Tips for figures and tables:

- *Number all figures (Figure 1, Figure 2...) and tables (Table 1, Table 2...)*
- *Include descriptive captions that can be understood without reading the main text*
- *Place figures/tables near where they're first referenced*
- *Ensure text in figures is legible*

5. Discussion and Limitations

What do your results mean? Discuss the broader implications for AI safety. What are the limitations of your work? What threat models or failure modes did you not address? What would you do differently?

Limitations

Be honest about constraints — methodological limitations, scope limitations, or aspects you couldn't fully address in the hackathon timeframe.

Future Work

What are the natural next steps? How could this work be extended?

6. Conclusion

Briefly summarize your main findings and their implications (1–2 paragraphs).

Code and Data

Include links if applicable. If your project doesn't involve code (e.g., policy analysis) or if there are info-hazard considerations, note that here.

- **Code repository:** *[Link to GitHub/GitLab if applicable]*
- **Data/Datasets:** *[Link if applicable]*
- **Other artifacts** (optional): *[Demo link, video walkthrough, Hugging Face Space, etc.]*

Author Contributions (optional)

[e.g., "A.B. led the project and designed experiments. C.D. implemented the code. All authors contributed to writing and reviewed the final manuscript."]

References

Use a consistent citation format. Include: Author(s), Year, Title, Venue/Publisher, and URL or DOI where available.

1. *[Reference 1]*
2. *[Reference 2]*
3. ...

Appendix (optional)

Supplementary material such as additional figures, detailed methodology, prompts used, extended results, etc.

LLM Usage Statement

If you used LLM assistance in developing your project or writing this report, briefly note how. Ensure all claims and results have been verified.

[e.g., "We used Claude to brainstorm approaches and help draft sections. All results and claims were independently verified."]